

dog_app

February 3, 2019

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: * Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/**/*.jpg"))
        dog_files = np.array(glob("/data/dog_images/**/*.jpg"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))

        # get bounding box for each detected face
        for (x,y,w,h) in faces:
            # add bounding box to color image
            cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)
```

```
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
```

```

img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
faces = face_cascade.detectMultiScale(gray)
return len(faces) > 0

```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

In [4]: `from tqdm import tqdm`

```

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

##-## Do NOT modify the code above this line. ##-##

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
print(f"Human faces detected in human images: {np.sum([face_detector(fn) for fn in human_files_short])}")
print(f"Human faces detected in dog images: {np.sum([face_detector(fn) for fn in dog_files_short])}")

```

Human faces detected in human images: 98%

Human faces detected in dog images: 17%

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

In [5]: *### (Optional)*
TODO: Test performance of another face detection algorithm.
Feel free to use as many code cells as needed.

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [6]: import torch
import torchvision.models as models

# define VGG16 model
VGG16 = models.vgg16(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
100%|| 553433881/553433881 [00:05<00:00, 99689972.36it/s]
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [7]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """
```

```

'''

## TODO: Complete the function.
## Load and pre-process an image from the given img_path
## Return the *index* of the predicted class for that image

img = cv2.imread(img_path)

img_trans = transforms.Compose(
    [transforms.ToPILImage(),
     transforms.Resize(256),
     transforms.CenterCrop(224),
     transforms.ToTensor()]
)

if use_cuda:
    img_tensor = img_trans(img).cuda()

VGG16.eval()

pred_class = VGG16(img_tensor.view(1,3,224,224)) # Re-size

return pred_class.argmax().item()

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [8]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.

    pred_idx = VGG16_predict(img_path)

    return 151 <= pred_idx <= 268

```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.

- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?

Answer:

```
In [9]: ### TODO: Test the performance of the dog_detector function
        ### on the images in human_files_short and dog_files_short.

        dogs_tp = np.sum([dog_detector(fn) for fn in dog_files_short])
        dogs_fn = np.sum([dog_detector(fn) for fn in human_files_short])

In [10]: print(f"{dogs_tp}% accuracy classifying dogs")
         print(f"{100-dogs_fn}% accuracy classifying humans as not dogs")

73% accuracy classifying dogs
100% accuracy classifying humans as not dogs
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [11]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [12]: import os
         from torchvision import datasets

         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

In [13]: # Dataloader parameters
         BATCH_SIZE = 30
         NUM_WORKERS=0

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         data_transform_train = transforms.Compose([transforms.RandomResizedCrop(260, scale=(0.7
                                                    transforms.RandomHorizontalFlip(p=.5),
                                                    transforms.RandomRotation(15),
                                                    transforms.CenterCrop(224),
                                                    transforms.ToTensor(),
                                                    transforms.Normalize([0.5, 0.5, 0.5], [0.5,
                                                    ])

         data_transform_test = transforms.Compose([transforms.Resize(260),
                                                    transforms.CenterCrop(224),
                                                    transforms.ToTensor(),
                                                    transforms.Normalize([0.5, 0.5, 0.5], [0.5,
                                                    ])

         data_categories = ['train', 'valid', 'test']
         loaders_scratch = {}
```



```

for cat in data_categories:

    data = datasets.ImageFolder(f"/data/dog_images/{cat}",
                                transform = (data_transform_train if cat == 'train' else
                                                data_transform_test))

    print(f"{len(data)} {cat} images")

    loaders_scratch[cat] = torch.utils.data.DataLoader(data, batch_size=BATCH_SIZE,
                                                         num_workers=NUM_WORKERS, shuffle=True)

6680 train images
835 valid images
836 test images

In [14]: print(iter(loaders_scratch['train']).next()[0].shape)
         iter(loaders_scratch['valid']).next()[0].shape

torch.Size([30, 3, 224, 224])

Out[14]: torch.Size([30, 3, 224, 224])

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: > - *RandomResizedCrop*: This randomly resizes & crops images to given size. I chose 260 to allow the transforms to have a larger canvas to work with before rotating. I also changed the default scale as I felt it went too low (no real basis other than my immature gut feel) > - *RandomHorizontalFlip*: Randomly flips image with $p=0.5$ - Standard way to augment > - *RandomRotation*: Up to 15 deg - Also standard > - *CenterCrop(224)*: This ultimate output size was chosen to match the VGG paper

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [15]: import torch.nn as nn
         import torch.nn.functional as F
         import torch
         import torchvision.models as models

         # define the CNN architecture
         class Net(nn.Module):
             ### TODO: choose an architecture, and complete the class
             def __init__(self, N):
                 super(Net, self).__init__()

```

```

# Input is an NxN image
self.N = N

## Define layers of a CNN
self.cv1 = nn.Conv2d(3, 64, 3, padding=1)
self.bn1 = nn.BatchNorm2d(64)

self.cv2 = nn.Conv2d(64, 128, 3, padding=1)
self.bn2 = nn.BatchNorm2d(128)

self.cv3 = nn.Conv2d(128, 256, 3, padding=1)
self.bn3 = nn.BatchNorm2d(256)
self.cv4 = nn.Conv2d(256, 256, 3, padding=1)
self.bn4 = nn.BatchNorm2d(256)

#self.cv5 = nn.Conv2d(256, 512, 3, padding=1)
#self.cv6 = nn.Conv2d(512, 512, 3, padding=1)

#self.cv7 = nn.Conv2d(512, 512, 3, padding=1)
#self.cv8 = nn.Conv2d(512, 512, 3, padding=1)

# Max pooling layer
self.mp = nn.MaxPool2d(2)

# Fully connected - N/4 * 128
self.fc1 = nn.Linear(self.N*self.N*4, 1024)
#self.fc2 = nn.Linear(4096, 4096)
self.fc3 = nn.Linear(1024, 133)

# Dropout
self.do = nn.Dropout(0.25)

def forward(self, x):
    ## Define forward behavior
    #print(f"Input {x.shape}")

    # N x N x 3 -> N/2 x N/2 x 64
    x = self.cv1(x)
    x = self.bn1(x)
    x = F.relu(x)
    x = self.mp(x)

    # N/2 x N/2 x 64 -> N/4 x N/4 x 128
    x = self.cv2(x)
    x = self.bn2(x)
    x = F.relu(x)
    x = self.mp(x)

```

```

# N/4 x N/4 x 128 -> N/8 x N/8 x 256
x = self.cv3(x)
x = self.bn3(x)
x = F.relu(x)
x = self.cv4(x)
x = self.bn4(x)
x = F.relu(x)
x = self.mp(x)

## N/8 x N/8 x 256 -> N/16 x N/16 x 512 = 2*N^2
#x = self.cv5(x)
#x = self.bn3(x)
#x = F.relu(x)
#x = self.cv6(x)
#x = self.bn4(x)
#x = F.relu(x)
#x = self.mp()

## N/16 x N/16 x 512 = 2*N^2
#x = self.cv7(x)
#x = self.bn3(x)
#x = F.relu(x)
#x = self.cv8(x)
#x = self.bn4(x)
#x = F.relu(x)
#x = self.mp()

# Flatten
x = x.view(-1, self.N*self.N*4)

x = self.do(x)
x = self.fc1(x)
x = F.relu(x)
x = self.do(x)
#x = self.fc2(x)
#x = F.relu(x)
#x = self.do(x)
x = self.fc3(x)

return x

#-#-# You do NOT have to modify the code below this line. #-#-#

# instantiate the CNN
model_scratch = Net(224)
print(model_scratch)
# move tensors to GPU if CUDA is available

```

```

if use_cuda:
    model_scratch.cuda()

```

```

Net(
  (cv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (cv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (cv3): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (cv4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (mp): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=200704, out_features=1024, bias=True)
  (fc3): Linear(in_features=1024, out_features=133, bias=True)
  (do): Dropout(p=0.25)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

Starting point is the VGG paper's A architecture: <https://www.kaggle.com/keras/vgg16/home>

Following decisions/ compromises were made: 1. Excluded last 4 conv layers & middle FC layer due to memory constraints 2. BatchNorm2d layers added - Suggestion to speed up from forums 3. The model started overfitting quite quickly so added dropout layers between fully connected layers

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a **loss function** and **optimizer**. Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [16]: import torch.optim as optim

```

```

In [17]: ### TODO: select loss function
         criterion_scratch = nn.CrossEntropyLoss()

```

```

### TODO: select optimizer
optimizer_scratch = optim.SGD(params=model_scratch.parameters(), lr=0.01)

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model_scratch.pt'.

```
In [18]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        print(f"Epoch: {epoch}")
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()

        for batch_idx, (data, target) in enumerate(loaders['train']):

            optimizer.zero_grad()

            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            prediction = model(data)
            loss = criterion(prediction, target)
            # Backward pass
            loss.backward()
            # Optimizer step
            optimizer.step()
            ## record the average training loss, using something like
            train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        #####
        # validate the model #
        #####

        print('Moving to Eval')

        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
```

```

        ## update the average validation loss
        prediction = model(data)
        loss = criterion(prediction,target)

        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss <= valid_loss_min:

        # REWRITE:
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.fo

        torch.save(model.state_dict(), save_path)

        valid_loss_min = valid_loss

# return trained model
return model

```

```

In [30]: %%time
# train the model
model_scratch = train(25, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

```

```

Epoch: 1
Moving to Eval
Epoch: 1      Training Loss: 4.738313      Validation Loss: 4.470997
Validation loss decreased (inf --> 4.470997). Saving model ...
Epoch: 2
Moving to Eval
Epoch: 2      Training Loss: 4.371082      Validation Loss: 4.271564
Validation loss decreased (4.470997 --> 4.271564). Saving model ...
Epoch: 3
Moving to Eval
Epoch: 3      Training Loss: 4.176295      Validation Loss: 4.159624
Validation loss decreased (4.271564 --> 4.159624). Saving model ...
Epoch: 4
Moving to Eval
Epoch: 4      Training Loss: 4.024529      Validation Loss: 4.087957
Validation loss decreased (4.159624 --> 4.087957). Saving model ...
Epoch: 5

```

```

Moving to Eval
Epoch: 5      Training Loss: 3.905927      Validation Loss: 4.020880
Validation loss decreased (4.087957 --> 4.020880). Saving model ...
Epoch: 6
Moving to Eval
Epoch: 6      Training Loss: 3.792347      Validation Loss: 3.870909
Validation loss decreased (4.020880 --> 3.870909). Saving model ...
Epoch: 7
Moving to Eval
Epoch: 7      Training Loss: 3.653920      Validation Loss: 3.969525
Epoch: 8
Moving to Eval
Epoch: 8      Training Loss: 3.593251      Validation Loss: 3.808012
Validation loss decreased (3.870909 --> 3.808012). Saving model ...
Epoch: 9
Moving to Eval
Epoch: 9      Training Loss: 3.502166      Validation Loss: 3.935073
Epoch: 10
Moving to Eval
Epoch: 10     Training Loss: 3.418550      Validation Loss: 3.679215
Validation loss decreased (3.808012 --> 3.679215). Saving model ...
Epoch: 11
Moving to Eval
Epoch: 11     Training Loss: 3.337316      Validation Loss: 3.660758
Validation loss decreased (3.679215 --> 3.660758). Saving model ...
Epoch: 12
Moving to Eval
Epoch: 12     Training Loss: 3.241210      Validation Loss: 3.710969
Epoch: 13
Moving to Eval
Epoch: 13     Training Loss: 3.146957      Validation Loss: 3.505728
Validation loss decreased (3.660758 --> 3.505728). Saving model ...
Epoch: 14
Moving to Eval
Epoch: 14     Training Loss: 3.069679      Validation Loss: 3.563920
Epoch: 15
Moving to Eval
Epoch: 15     Training Loss: 2.999975      Validation Loss: 4.005156
Epoch: 16
Moving to Eval
Epoch: 16     Training Loss: 2.916625      Validation Loss: 3.694385
Epoch: 17
Moving to Eval
Epoch: 17     Training Loss: 2.841942      Validation Loss: 3.502520
Validation loss decreased (3.505728 --> 3.502520). Saving model ...
Epoch: 18
Moving to Eval
Epoch: 18     Training Loss: 2.750790      Validation Loss: 3.476121

```

```

Validation loss decreased (3.502520 --> 3.476121). Saving model ...
Epoch: 19
Moving to Eval
Epoch: 19      Training Loss: 2.685582      Validation Loss: 3.606239
Epoch: 20
Moving to Eval
Epoch: 20      Training Loss: 2.610869      Validation Loss: 3.512245
Epoch: 21
Moving to Eval
Epoch: 21      Training Loss: 2.513334      Validation Loss: 3.617355
Epoch: 22
Moving to Eval
Epoch: 22      Training Loss: 2.455744      Validation Loss: 3.628527
Epoch: 23
Moving to Eval
Epoch: 23      Training Loss: 2.363952      Validation Loss: 3.473887
Validation loss decreased (3.476121 --> 3.473887). Saving model ...
Epoch: 24
Moving to Eval
Epoch: 24      Training Loss: 2.328944      Validation Loss: 3.535422
Epoch: 25
Moving to Eval
Epoch: 25      Training Loss: 2.214098      Validation Loss: 3.487274
CPU times: user 58min 53s, sys: 5min 5s, total: 1h 3min 59s
Wall time: 57min 45s

```

```

In [19]: # load the model that got the best validation accuracy
         model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [20]: def test(loaders, model, criterion, use_cuda):

         # monitor test loss and accuracy
         test_loss = 0.
         correct = 0.
         total = 0.

         model.eval()
         for batch_idx, (data, target) in enumerate(loaders['test']):
             # move to GPU
             if use_cuda:
                 data, target = data.cuda(), target.cuda()
             # forward pass: compute predicted outputs by passing inputs to the model

```



```

        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

```

```
In [21]: # call test function
```

```
        test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 3.477499

Test Accuracy: 22% (187/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [22]: # Dataloader parameters
```

```
BATCH_SIZE = 30
```

```
NUM_WORKERS=0
```

```
### TODO: Write data loaders for training, validation, and test sets
```

```
## Specify appropriate transforms, and batch_sizes
```

```
data_transform_train = transforms.Compose([transforms.RandomResizedCrop(260, scale=(0.7
                                                    transforms.RandomHorizontalFlip(p=.5),
                                                    transforms.RandomRotation(15),
```

```

        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.5, 0.5, 0.5], [0.5,
])

data_transform_test = transforms.Compose([transforms.Resize(260),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.5, 0.5, 0.5], [0.5,
])

data_categories = ['train', 'valid', 'test']
data_transfer = {}
loaders_transfer = {}

for cat in data_categories:

    data_transfer[cat] = datasets.ImageFolder(f"/data/dog_images/{cat}",
        transform = (data_transform_train if cat == 'train' else

    print(f"{len(data)} {cat} images")

    loaders_transfer[cat] = torch.utils.data.DataLoader(data_transfer[cat], batch_size=
        num_workers=NUM_WORKERS, shuffle=True)

836 train images
836 valid images
836 test images

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [23]: import torch.nn as nn
import torch.nn.functional as F
import torch
import torchvision.models as models

model_transfer = models.vgg16(pretrained=True)

print(model_transfer)

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)

```

```

(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(3): ReLU(inplace)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(6): ReLU(inplace)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)

```

We need to replace the classifier with one out size 133

```
In [24]: model_transfer.classifier[6] = nn.Linear(in_features=4096, out_features=133, bias=True)
```

Freeze training on features:

```
In [25]: for param in model_transfer.features.parameters():
          param.requires_grad=False
```

```
In [26]: if use_cuda:
        model_transfer = model_transfer.cuda()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

Steps taken: 1. Fetch pretrained model: VGG16, as we saw in a previous lesson 2. Change just the output layer to predict 133 dog breeds

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [27]: criterion_transfer = nn.CrossEntropyLoss()

        # Optimise just the classifier parameters:
        optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [ ]: %%time

        # train the model
        model_transfer = train(50, loaders_transfer, model_transfer,
                                optimizer_transfer, criterion_transfer, use_cuda, 'model_transfer.pt')

Epoch: 1
Moving to Eval
Epoch: 1      Training Loss: 4.445627      Validation Loss: 3.628706
Validation loss decreased (inf --> 3.628706).  Saving model ...
Epoch: 2
Moving to Eval
Epoch: 2      Training Loss: 3.139152      Validation Loss: 2.176020
Validation loss decreased (3.628706 --> 2.176020).  Saving model ...
Epoch: 3
Moving to Eval
Epoch: 3      Training Loss: 2.048873      Validation Loss: 1.343856
Validation loss decreased (2.176020 --> 1.343856).  Saving model ...
Epoch: 4
Moving to Eval
Epoch: 4      Training Loss: 1.516557      Validation Loss: 0.999414
Validation loss decreased (1.343856 --> 0.999414).  Saving model ...
Epoch: 5
Moving to Eval
Epoch: 5      Training Loss: 1.252584      Validation Loss: 0.842267
```

Validation loss decreased (0.999414 --> 0.842267). Saving model ...
Epoch: 6
Moving to Eval
Epoch: 6 Training Loss: 1.092260 Validation Loss: 0.756083
Validation loss decreased (0.842267 --> 0.756083). Saving model ...
Epoch: 7
Moving to Eval
Epoch: 7 Training Loss: 0.987736 Validation Loss: 0.695668
Validation loss decreased (0.756083 --> 0.695668). Saving model ...
Epoch: 8
Moving to Eval
Epoch: 8 Training Loss: 0.920545 Validation Loss: 0.656052
Validation loss decreased (0.695668 --> 0.656052). Saving model ...
Epoch: 9
Moving to Eval
Epoch: 9 Training Loss: 0.844562 Validation Loss: 0.620181
Validation loss decreased (0.656052 --> 0.620181). Saving model ...
Epoch: 10
Moving to Eval
Epoch: 10 Training Loss: 0.799041 Validation Loss: 0.599406
Validation loss decreased (0.620181 --> 0.599406). Saving model ...
Epoch: 11
Moving to Eval
Epoch: 11 Training Loss: 0.786013 Validation Loss: 0.576641
Validation loss decreased (0.599406 --> 0.576641). Saving model ...
Epoch: 12
Moving to Eval
Epoch: 12 Training Loss: 0.738657 Validation Loss: 0.572876
Validation loss decreased (0.576641 --> 0.572876). Saving model ...
Epoch: 13
Moving to Eval
Epoch: 13 Training Loss: 0.706403 Validation Loss: 0.557193
Validation loss decreased (0.572876 --> 0.557193). Saving model ...
Epoch: 14
Moving to Eval
Epoch: 14 Training Loss: 0.658957 Validation Loss: 0.547270
Validation loss decreased (0.557193 --> 0.547270). Saving model ...
Epoch: 15
Moving to Eval
Epoch: 15 Training Loss: 0.659556 Validation Loss: 0.535066
Validation loss decreased (0.547270 --> 0.535066). Saving model ...
Epoch: 16
Moving to Eval
Epoch: 16 Training Loss: 0.649434 Validation Loss: 0.526287
Validation loss decreased (0.535066 --> 0.526287). Saving model ...
Epoch: 17
Moving to Eval
Epoch: 17 Training Loss: 0.620054 Validation Loss: 0.522992

Validation loss decreased (0.526287 --> 0.522992). Saving model ...
Epoch: 18
Moving to Eval
Epoch: 18 Training Loss: 0.613136 Validation Loss: 0.513782
Validation loss decreased (0.522992 --> 0.513782). Saving model ...
Epoch: 19
Moving to Eval
Epoch: 19 Training Loss: 0.571425 Validation Loss: 0.504678
Validation loss decreased (0.513782 --> 0.504678). Saving model ...
Epoch: 20
Moving to Eval
Epoch: 20 Training Loss: 0.587055 Validation Loss: 0.504217
Validation loss decreased (0.504678 --> 0.504217). Saving model ...
Epoch: 21
Moving to Eval
Epoch: 21 Training Loss: 0.562756 Validation Loss: 0.502751
Validation loss decreased (0.504217 --> 0.502751). Saving model ...
Epoch: 22
Moving to Eval
Epoch: 22 Training Loss: 0.551554 Validation Loss: 0.497697
Validation loss decreased (0.502751 --> 0.497697). Saving model ...
Epoch: 23
Moving to Eval
Epoch: 23 Training Loss: 0.527800 Validation Loss: 0.501619
Epoch: 24
Moving to Eval
Epoch: 24 Training Loss: 0.515090 Validation Loss: 0.484503
Validation loss decreased (0.497697 --> 0.484503). Saving model ...
Epoch: 25
Moving to Eval
Epoch: 25 Training Loss: 0.511221 Validation Loss: 0.479480
Validation loss decreased (0.484503 --> 0.479480). Saving model ...
Epoch: 26
Moving to Eval
Epoch: 26 Training Loss: 0.488935 Validation Loss: 0.484479
Epoch: 27
Moving to Eval
Epoch: 27 Training Loss: 0.474668 Validation Loss: 0.472875
Validation loss decreased (0.479480 --> 0.472875). Saving model ...
Epoch: 28
Moving to Eval
Epoch: 28 Training Loss: 0.490038 Validation Loss: 0.475958
Epoch: 29
Moving to Eval
Epoch: 29 Training Loss: 0.484024 Validation Loss: 0.469292
Validation loss decreased (0.472875 --> 0.469292). Saving model ...
Epoch: 30
Moving to Eval

```

Epoch: 30          Training Loss: 0.446966          Validation Loss: 0.480300
Epoch: 31
Moving to Eval
Epoch: 31          Training Loss: 0.449989          Validation Loss: 0.467925
Validation loss decreased (0.469292 --> 0.467925). Saving model ...
Epoch: 32
Moving to Eval
Epoch: 32          Training Loss: 0.450739          Validation Loss: 0.472817
Epoch: 33
Moving to Eval
Epoch: 33          Training Loss: 0.441568          Validation Loss: 0.473826
Epoch: 34
Moving to Eval
Epoch: 34          Training Loss: 0.432885          Validation Loss: 0.468006
Epoch: 35
Moving to Eval
Epoch: 35          Training Loss: 0.417005          Validation Loss: 0.465467
Validation loss decreased (0.467925 --> 0.465467). Saving model ...
Epoch: 36
Moving to Eval
Epoch: 36          Training Loss: 0.427821          Validation Loss: 0.467103
Epoch: 37
Moving to Eval
Epoch: 37          Training Loss: 0.420244          Validation Loss: 0.466600
Epoch: 38
Moving to Eval
Epoch: 38          Training Loss: 0.403681          Validation Loss: 0.463843
Validation loss decreased (0.465467 --> 0.463843). Saving model ...
Epoch: 39
Moving to Eval
Epoch: 39          Training Loss: 0.406210          Validation Loss: 0.463697
Validation loss decreased (0.463843 --> 0.463697). Saving model ...
Epoch: 40
Moving to Eval
Epoch: 40          Training Loss: 0.386933          Validation Loss: 0.473116
Epoch: 41
Moving to Eval
Epoch: 41          Training Loss: 0.393897          Validation Loss: 0.461776
Validation loss decreased (0.463697 --> 0.461776). Saving model ...
Epoch: 42
Moving to Eval
Epoch: 42          Training Loss: 0.375833          Validation Loss: 0.454575
Validation loss decreased (0.461776 --> 0.454575). Saving model ...
Epoch: 43
Moving to Eval
Epoch: 43          Training Loss: 0.365521          Validation Loss: 0.466558
Epoch: 44
Moving to Eval

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [33]: # load the model that got the best validation accuracy (uncomment the line below)
         model_transfer.load_state_dict(torch.load('model_transfer.pt'))

         test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.509282

Test Accuracy: 83% (701/836)

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [29]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

         # list of class names by index, i.e. a name can be accessed like class_names[0]
         class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]

         def predict_breed_transfer(img_path):
             # load the image and return the predicted breed

             # Read in img using OpenCV
             img_array = cv2.imread(img_path)

             # Need to convert to PIL before using previous transforms
             img_PIL = transforms.ToPILImage()(img_array)

             img_tensor = data_transform_test(img_PIL)

             # Into required dimensions
             img_tensor = img_tensor.view(1,3,224,224)

             # To GPU
             if use_cuda:
                 img_tensor = img_tensor.cuda()

             model_transfer.eval()
```




Sample Human Output

```
pred_class = class_names[model_transfer(img_tensor).argmax().item()]

return pred_class
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [34]: ### TODO: Write your algorithm.
         ### Feel free to use as many code cells as needed.

def run_app(img_path, noisy=True):
    ## handle cases for a human face, dog, and neither

    dog_detected, human_detected = dog_detector(img_path), face_detector(img_path)

    if dog_detected or human_detected:
        breed = predict_breed_transfer(img_path)

        if dog_detected:
            print("Detected a dog")
            print(f"Predicted breed: {breed}")
        else:
            print("Detected a human")
            print(f"Closest dog breed resembled: {breed}")
```

```

        return breed
    else:
        print("Nothing detected")

```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement) Human & dogs are correctly separated. Only 2/5 dog breeds correctly identified in the sample below- this is worse than the expected %age from the test above.

1. Used the same normalisation of images for the dog breed detector as originally specified (means & s.d's all 0.5). The VGG model expects a different normalisation for each colour channel - this could provide an improvement.
2. In the model training, 50 epochs were allowed, but the operation timed out after 44, for which we saw an improvment in testing accraucy in the epoch previous. This suggests that the model should be run to at least 50 epochs, probably more. Could use early stopping here.
3. The algorithm currently only outputs the best match of dog. It could be altered to also produce the %age likelihood of the top 3

```

In [31]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.

```

```

In [35]: ## suggested code, below
import random

```

```

for file in np.hstack((random.choices(human_files, k=5), random.choices(dog_files, k=5))):
    print(file)
    run_app(file)
    print(" ---- ")

```

```

/data/lfw/Garry_Trudeau/Garry_Trudeau_0002.jpg

```

```

Detected a human

```

```

Closest dog breed resembled: Dachshund

```

```

----

```

```

/data/lfw/Martin_Short/Martin_Short_0001.jpg

```

```

Detected a human

```

Closest dog breed resembled: Dachshund

/data/lfw/Penelope_Ann_Miller/Penelope_Ann_Miller_0002.jpg
Detected a human
Closest dog breed resembled: Chinese crested

/data/lfw/Teri_ORourke/Teri_ORourke_0001.jpg
Detected a human
Closest dog breed resembled: Dogue de bordeaux

/data/lfw/Norah_Jones/Norah_Jones_0015.jpg
Detected a human
Closest dog breed resembled: Chinese crested

/data/dog_images/train/028.Bluetick_coonhound/Bluetick_coonhound_01985.jpg
Detected a dog
Predicted breed: German shorthaired pointer

/data/dog_images/valid/015.Basset_hound/Basset_hound_01109.jpg
Detected a dog
Predicted breed: Plott

/data/dog_images/train/053.Cocker_spaniel/Cocker_spaniel_03777.jpg
Detected a dog
Predicted breed: Cocker spaniel

/data/dog_images/test/054.Collie/Collie_03835.jpg
Detected a dog
Predicted breed: Alaskan malamute

/data/dog_images/train/096.Labrador_retriever/Labrador_retriever_06446.jpg
Detected a dog
Predicted breed: Labrador retriever
